



Containers in Research

Hands-on introduction to leveraging containers in research

Luis Gracia, PhD

 [lgracia](#)

 [luisico](#)

July 2024


Agenda

- What to expect from this workshop
- Introduction to containers
- Why should we care in scientific research
- How to get started – Hands on
 - Running containers
 - Creating container images
 - Orchestrating containers

What are Containers?

*"A container is a standard unit of software that **packages up code and all its dependencies** so the application runs quickly and reliably from one computing environment to another. A Docker container image is a lightweight, standalone, executable package of software that includes everything needed to run an application: code, runtime, system tools, system libraries and settings." – docker.com*

- Fast startup, low resources
- "Isolated" running context
- Can be further configured at runtime
- Standard (OCI)



credits: docker.com

→ Packaged as an Image, runs as a Container

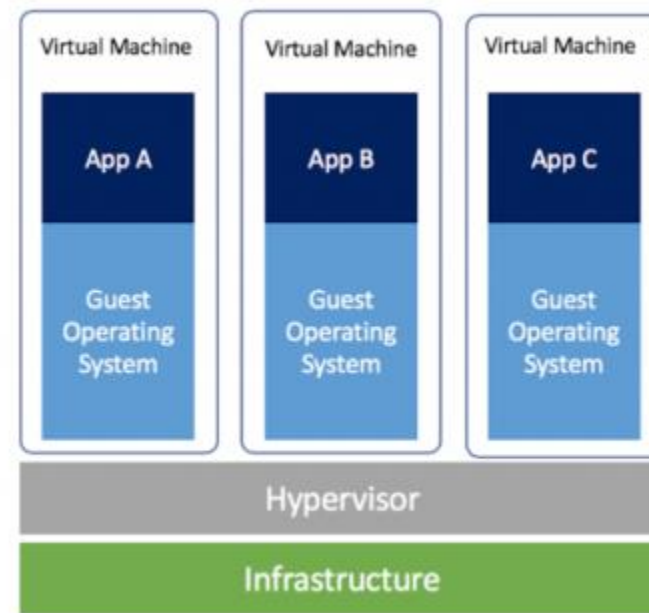
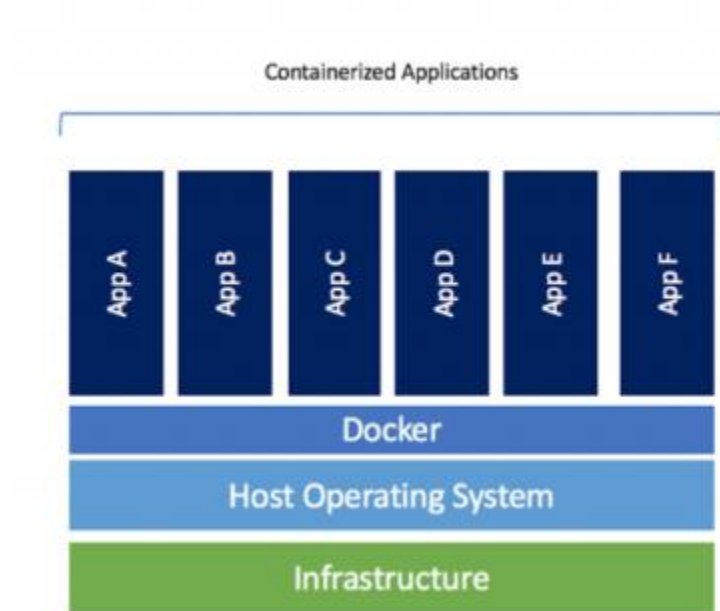
→ Depends on the OS kernel:

- **Linux**
- Mac
- Windows

Why should we care in research?

- Reproducibility
 - App and dependencies are pre-installed and configured in the image
 - Defined in code (track, version)
 - Use versioning to run old workflows years later
 - Hardware and kernel (OS) changes can affect reproducibility
- Portability and collaboration
 - Compile/Package once, run multiple times
 - Run same image in different systems:
 - laptop, HPC, cloud
 - Development, test, production
 - Share software or pipelines with colleagues
 - Use code from lab mates
 - Other labs can test/review your code/results
 - Run standard software or develop your own
- Isolated and conflict-free environment
 - Avoid “dependency hell” and conflicts between libraries, languages and dependencies
 - Use different versions of python simultaneously
 - Test different versions of a library
 - Run different environments simultaneously (i.e. tests parameters)
- Near native operation
 - Low computing resources (memory, cpu)
 - Fast start
 - Run multiple containers simultaneously
- Alternatives:
 - Self managed
 - Software repositories
 - Distro packages (deb, rpm)
 - Distro-agnostic packages: [Flatpack](#), [Snaps](#), [AppImage](#)
 - Version managers: [nvm](#), [nodenv](#), [Renv](#), [rbenv](#)
 - Language specific:
 - General: [asdf](#), [mise](#), [anyenv](#)
 - VMs

Containers vs VMs



credits: docker.com

Very fast	Startup	Slow
Low mem/cpu	Resource usage	High mem/cpu (reserved)
Small (lots of tricks to reduce it)	Size	Big
Yes, but can get tricky	Isolation	Yes
App and dependencies; runs on same OS	Portability/Maintenance	OS, app and dependencies; runs anywhere
Yes	Versioning	Yes
Easy	Provenance	Difficult
Easy	Orchestration	Difficult

Hands-on Instructions

- How to follow along with Docker
 - Use provided VMs at <https://container-workshop.cloud.edu.au/guacamole>
 - One pre-allocated VM per person, login with your credentials
 - ❖ Username: the complete name you registered with
 - ❖ Password: email address you registered with
 - Web based terminals
 - Instructors can access your terminal to help troubleshooting (pair programming style)
 - No GUI
 - No direct access to VM via SSH or browser
 - Access configuration with CTRL+SHIFT+ALT (or CTRL+SHIFT+OPT on Macs)
 - Docker installed on your own computer (<https://docs.docker.com/get-docker>)
 - [Play with Docker](#) online
- Content at <https://github.com/UoA-eResearch/resbaz-2024-containers>
- Need help?
 - We assume you have some Linux shell terminal knowledge
 - Contact us (Luis, Andre, Jason) via Zoom chat
 - Feel free to contact me (<https://profiles.auckland.ac.nz/luis-gracia-valen>) via email after the session
- Docker alternatives
 - [Podman](#)
 - [Apptainer](#) (Singularity)
 - [LXC](#)

Hands-on: some examples

- Hello world: sum "2 + 3"
 - Python
 - Node.js
 - R
- Interactive R
- Serving websites with Nginx
- Local RStudio via browser

Docker main commands

Reference: <https://docs.docker.com/reference/cli/docker>

```
docker
  run      # Create and run a new container from an image
  ps       # List containers
  rm       # Remove one or more containers
  kill     # Kill one or more running containers
  exec     # Execute a command in a running container
  pull     # Download an image from a registry
  push     # Upload an image to a registry
  image    # Operations with images
  cp       # Copy files/folders between a container and the local filesystem
  logs     # Fetch the logs of a container
```


Running Containers

```
docker run [options] image [command]
```

<code>--rm</code>			# Remove the container when it exits
<code>--interactive</code>	<code>-i</code>		# Interactive session
<code>--tty</code>	<code>-t</code>		# Allocate a pseudo-TTY
<code>--name</code>		name	# Assign a name to the container
<code>--detach</code>	<code>-d</code>		# Run container in background
<code>--volume</code>	<code>-v</code>	host_dir:container_dir	# Share a host directory inside the container
<code>--workdir</code>	<code>-w</code>	container_dir	# Starting directory inside the container
<code>--port</code>	<code>-p</code>	host_port:container_port	# Publish a container's port(s) to the host
<code>--env</code>	<code>-e</code>	VAR=value	# Set environment variable inside container
<code>--env-file</code>		/path/to/file	# Pass environment variables from file
<code>--user</code>	<code>-u</code>	user	# Use username in container
<code>--entrypoint</code>		/path/to/executable	# Overwrite the default entrypoint of the image

Hands-on: running containers

- Running, listing and cleaning containers
- Managing container names
- Executing non-default commands
- Running interactive containers
- Running containers in the background
- Exposing container ports to the host machine (or the internet)
- Configuring containers: sharing files and variables
- Configuring commands ("entrypoints")
- Dealing with permissions when sharing files

Container Images

- A container is an instance of an image
- When a container stops, changes are lost (i.e. pip install)
- Naming convention: `registry/org/name:tag`
 - `registry:`
 - Docker Hub: <https://hub.docker.com> (default)
 - Quay: <https://quay.io/search>
 - Git platform: GitHub, GitLab
 - Cloud: AWS, Google Cloud, Azure
 - Private
 - `org/name:`
 - Usually called "repository" or image name
 - Only `name` is mandatory
 - `tag:`
 - Default is "latest"
 - Arbitrary text, no special meaning

Examples:

- `ubuntu`
- `ubuntu:24.04`
- `ubuntu/mysql`
- `quay.io/quay/ubuntu`
- `localhost:5000/testing/ubuntu:20240709`

Images main commands

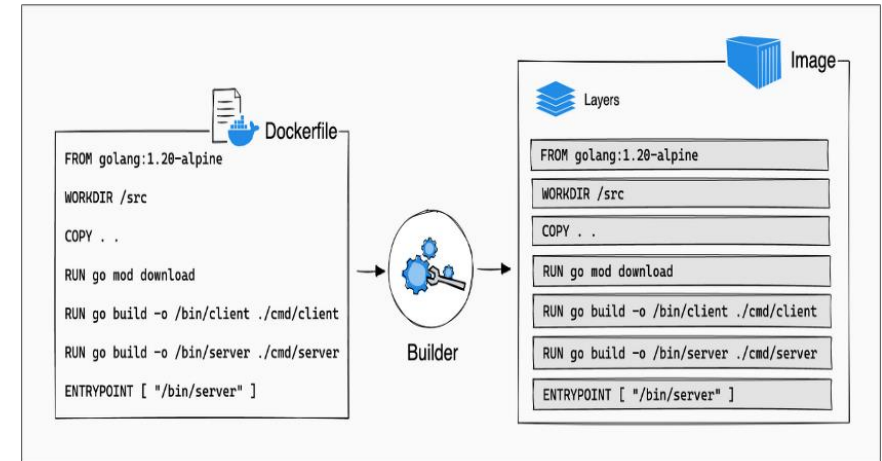
```
docker image
  ls          # List images (or `docker images`)
  pull        # Download an image from a registry
  push        # Upload an image to a registry
  rm          # Remove an image
  prune       # Remove unused images
  build       # Build an image from a Dockerfile (`docker build`)
  inspect     # Inspect properties of an image (`docker inspect`)
  history     # Show the build history of an image
  save        # Export an image as a tar archive
  load        # Import a tarball as an image
```

Pulling (downloading) Images

- `docker pull image`
 - Image is downloaded layer by layer, existing layers being reused
 - `docker run` will pull the image if not present locally, but it will not check if there is a newer version in the registry
 - `docker pull` will download newer versions of an image already present locally
- Inspect: `docker inspect image`
- Keep an eye on disk space: `docker system df`

Building Container Images

- Dockerfile: define how the image is built using code:
 - Selecting a base image to start with
 - Adding, removing and changing files
 - Running commands
 - Defining ports, volumes, variables, commands, etc...
 - Reference: <https://docs.docker.com/reference/dockerfile>
- Layers:
 - New layer every time the filesystem is changed (i.e. adding files, running a command that changes files)
 - Layers are downloaded independently and composed on top of each other to create the final image
 - Shared among images, reducing disk space
- Metadata:
 - Labels (authorship, copyright, versions)
 - Ports
 - Volumes



credits: docker.com

```
docker build [options] path
--tag      -t    Tag name           # Image name
--file     -f    Dockerfile         # Path to Dockerfile
--progress auto|plain|tty          # Verbosity
--build-arg ARG=value              # Pass build argument
--no-cache
```


Hands-on: Images

- Finding images in Docker Hub
- Pulling images locally
- Creating images from a Dockerfile
 - Develop and share own software (*Campsites*)
 - Package open source software (*Fastutils*)
 - Compile open source software (*Open Babel*)
- Sharing images
 - Push to registry (need an account in [Docker Hub](#) or another registry)
 - Save/Load

Hands-on: DOC Campsites

Dockerfile

```
FROM python:3.12.4-alpine3.20
WORKDIR /app
ADD requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
ADD . .
CMD ["python", "tally.py"]
```

requirements.txt

```
pandas==2.2.2
```

data.csv

```
Region,Number of powered sites,Number of unpowered sites
Central North Island,0,12
Fiordland,0,5
```

tally.py

```
data = pd.read_csv('data.csv', header = 0, usecols =
                    ['Region',
                     'Number of powered sites',
                     'Number of unpowered sites'
                    ])

print("Campsites by region")
print("-----")
campsites_by_region = data.groupby('Region').Region.count()
print(campsites_by_region)
print("")
print("Site types per region")
print("-----")
site_types_per_region = data.groupby('Region').sum()
print(site_types_per_region)
```

Hands-on: Sharing Images

- Push to registry
- Create account in [Docker Hub](#)
 1. Login into your account from your machine
 2. Tag image to include your `org` (your account name)
 3. Push (upload) your image

```
docker login
docker tag campites:1.0.0 org/campsites:1.0.0
docker push org/campsites:1.0.0
```

- Save / Load (manual process)

```
docker image save campsites:1.0.0 -o campsites-v1.0.0.tar
```

Transfer (ssh, rsync, shared drive, etc)

```
docker image load -i campsites-v1.0.0.tar
```

Hands-on: Fastutils

Dockerfile

```
FROM debian:bookworm-slim

RUN apt-get update \

    && apt-get install -y git make g++ zlib1g-dev \

    && rm -rf /var/lib/apt/lists/*

RUN git clone https://github.com/haghshenas/fastutils.git /tmp/fastutils \

    && cd /tmp/fastutils \

    && make \

    && cp fastutils /usr/local/bin \

    && rm -rf /tmp/fastutils

WORKDIR /app

CMD ["/usr/local/bin/fastutils"]
```

Dockerfile.alpine

```
FROM alpine:3.20 AS builder

RUN apk --no-cache add \

    g++ \

    git \

    make \

    zlib-dev \

    zlib-static

RUN git clone https://github.com/haghshenas/fastutils.git /tmp/fastutils \

    && cd /tmp/fastutils \

    && env LDFLAGS=-static make

FROM alpine:3.20

COPY --from=builder /tmp/fastutils /usr/local/bin/

WORKDIR /app

CMD ["/usr/local/bin/fastutils"]
```

Hands-on: Open Babel

Dockerfile

```
FROM alpine:3.20
ARG OPENBABEL_REF=master
RUN apk --no-cache add \
    git \
    g++ cmake make perl \
    boost-dev cairo-dev eigen-dev libxml2-dev zlib-dev
RUN git clone https://github.com/openbabel/openbabel.git /tmp/openbabel \
    && cd /tmp/openbabel \
    && git checkout ${OPENBABEL_REF} \
    && env CXXFLAGS="-Wno-deprecated-declarations" cmake . \
    && make -j2 \
    && make install \
    && rm -rf /tmp/openbabel
ENTRYPOINT ["/usr/local/bin/obabel"]
CMD [""]
```

Orchestrating Containers

Docker Compose:

- Declarative instructions to run/orchestrate multiple images with custom configuration
- Reference: <https://docs.docker.com/compose/compose-file/>
- Per directory

docker-compose.yml

```
services:
  app:
    image: best-app:latest
    environment:
      VAR1: value1
    ports:
      - '8000:80'
    depends_on:
      - database
  database:
    image: mysql
    environment:
      MYSQL_ROOT_PASSWORD: secure
```

docker compose

up	# Create and start services
down	# Stop and remove services
ps	# List containers
exec	# Execute command in container
pull	# Pull (download) images
logs	# View output from containers
build	# Build images

Hands-on: RStudio

Setup RStudio with docker compose and pre-installed packages

docker-compose.yml

```
services:
  rstudio:
    image: rocker/verse:4.4
    restart: unless-stopped
    environment:
      PASSWORD: rstudio
    ports:
      - '8787:8787'
    volumes:
      - ./home/rstudio
```

docker-compose.yml

```
services:
  rstudio:
    image: my-rstudio
    build: .
    restart: unless-stopped
    environment:
      PASSWORD:
    ports:
      - '8787:8787'
    volumes:
      - ./home/rstudio
```

Dockerfile

```
FROM rocker/verse:4.4
RUN install2.r --error --skipinstalled \
  janitor \
  && rm -rf /tmp/downloaded_packages
```